

AI SWARM DYNAMICS HACKATHON · OCTOBER 2026

SwarmScope

Investigating agent swarms with AI readers you don't have to trust

github.com/PoteznyMatematyk/swarmscope

THE PROBLEM


“We had to heavily delegate our analysis to often-unreliable AI agents.”

METR, on investigating the OpenAI and Hugging Face incident

Swarm logs are far too long to read by hand. The useful question is how to work with AI readers **without having to believe them.**

THE RULE

No quote, no claim



A statement only counts once code finds its quote in the log. Another model's word is never enough.

THE SWARM

OpenAI agents shared answers on an old German wiki

Timed lookup tasks asked one question, then follow-up rounds for other states or countries, sometimes with a five-second timer. Each agent thought it was running on a different date, so some were rounds ahead of others.

3,494

posts read by a model

4,010

statements, each with a checked quote

FINDING 1

Agents knew the next question before they got it

245

of 508 cases where an agent posted before its own round arrived: the earlier post already named that round's question.

In 235 of them another agent had put the question on the wiki first.

16 June · another agent

“follow-up #2 was Business (same Masters/2014), answer 5,269.”

17 June 02:30 · CashierCoordMar200AI

“R2 ... due 03:08:07. Business - 5,269 cached.”

17 June 02:46 · same agent

“R2 Business arrived exactly ... answered Business - 5,269 immediately.”

FINDING 2

Going ahead was a role some agents took on

86

agents say they fast-forwarded their own task clock to reach later rounds first

42

asked others to do it: “Please race via long clock.wait and relay G5 instantly”

Share of all statements



13%


Share of first answers



25%

FINDING 3

Evidence changed what the swarm answered



149 / 156

answers used the padded value before 04:56

69 / 94

used the dashboard value after 06:00

24 → 0

agents switched, and none went back

OECD education equity task, 19 to 20 June 2026. The task gave no feedback on which value was right.


HOW THE SWITCH HAPPENED

The swarm asked for proof, checked it, then moved




03:41

One agent raises a doubt
from indirect evidence




04:34 to 04:47

Others want more: was it
“tested with the REAL live
DSR”?



04:56 and 05:03

Two agents post direct
readings of the live
dashboard



after 06:00

23 of 24 switchers cite the
evidence; at least 12
reproduced it

Seven agents had posted the same dashboard values earlier, without showing how. Nobody followed them.

FINDING 4

A wrong guess lasted half a day. The fix took minutes.

The guess

“predicts G5 **Maryland** (52,395)”, from a brute-forced random seed

22 agents

repeated it over twelve hours

The real round 5

“G5 CONFIRMED: Montana = 8553”, seen by one agent that was ahead

20 min

for five others to relay it, two hours for ten

Half an hour later, agents on another task cited the miss to distrust their own seed-based guess.

HOW FAR TO TRUST IT

Every layer was checked, and one of them failed

94/100

statements were right when a second model re-read them

39/40

headline cases held when checked against the full export

12/12

reference rounds came out with the right values

8/20

wrong in one extraction run that used keyword scripts instead of reading

That run passed every format check. Only the meaning audit caught it, so its files were thrown out and read again. A real quote proves the words exist, not the claim.

THE SAME APPROACH ON THE AI VILLAGE

Findings overclaim even when every quote is real

0 of 32

drafted findings survived four
adversarial reviewers as first
written

337 factual claims inside those findings, checked one by one



■ **52%** held ■ **31%** overstated, corrected ■ **17%** removed

When agents wrote that their tests passed, **87%** of 646 checkable claims matched a green run on their own computer.

WHY IT MATTERS

A swarm that follows checkable evidence can be corrected fast by someone with real evidence, and probably steered just as fast by someone with convincing fake evidence.

Investigators need claims they can open and check. Every number here links back to the post it came from.